

Thomas Nguyễn

0901463220 | thomasnguyen10092002@gmail.com | [GITHUB](#) | [LINKEDIN](#) | [PORTFOLIO](#)

PROFESSIONAL SUMMARY

- AI Engineer specializing in Computer Vision, Deep Learning, and MLOps pipelines.
- Strong proficiency in Python and PyTorch, combined with practical experience in automation testing framework engineering and advanced Generative AI applications (RAG & Multi-Agent systems)

TECHNICAL SKILLS

- AI & Generative AI:** PyTorch, OpenCV, LlamaIndex, CrewAI, Mediapipe, Scikit-learn, LangGraph
- Automation & MLOps:** Ansible (YAML), Jenkins, DVC, MLflow, Airflow, Docker, Docker Compose, GitHub Actions.
- Backend & Cloud:** Python (FastAPI, Flask), Celery, Redis, PostgreSQL, Azure, Linux.
- Monitoring & Tools:** Prometheus, Grafana, Jira, Bitbucket, Git.

EDUCATION

Ton Duc Thang university (TDTU)

Bachelor of Computer Science - GPA: 7.3/10

Ho Chi Minh City, Viet Nam

September 2020 - September 2025

Coursework: Object-Oriented Programming, Data Structure and Algorithms, Database Management Systems, Machine Learning, Computer Vision, Natural Language Processing

Honors:

- Second prize for students researching integration with Computer Vision at Ton Duc Thang University (2021 - 2022)
- Second Prize in the competition on correcting IPA pronunciation for Vietnamese learners (2024–2025)

WORK EXPERIENCES

Software Engineer 1

Terralogic Company

April 2025 - Present

Ho Chi Minh City, Viet Nam

Core Infrastructure & Automation (Cisco Enterprise Project):

- Designed, optimized, and maintained complex **Ansible playbooks (YAML)** to automate system configuration and functional validation.
- Integrated automated Regression Test suites into **Jenkins CI/CD pipelines**, leveraging **Bitbucket** and **Jira** to accelerate delivery cycles and sustain high infrastructure stability.

AI R&D Initiative (Speech-to-IPA Product):

- Awarded Second Prize in the company's internal AI development competition for independently pioneering an end-to-end Speech-to-IPA pronunciation correction system.
- Fine-tuned a **deep learning** model using **PyTorch**, integrating it with a **Flask** backend and **Next.js** frontend to build the functional product.
- Optimized local inference to achieve low-latency execution with high precision, rigorously evaluating model performance using **Word Error Rate (WER)** and **Character Error Rate (CER)** metrics.

AI Engineer Intern

TMA Solution Lab 6

August 2024 - November 2024

Ho Chi Minh City, Viet Nam

- Developed a sign language recognition model, achieving **94% accuracy** using **CNN1D + Transformer**.
- Preprocessed data using **Mediapipe & OpenCV**, extracting 543 landmark points and applying **Z-score normalization**.
- Optimized model performance through **K-fold cross-validation**, **Focal Loss**, and **Affine Data Augmentation**.
- Built **Flask APIs** and converted **PyTorch** models to **ONNX** format for optimized inference.
- Evaluated model effectiveness using the **Confusion Matrix**, **F1-score**, **Precision**, and **Recall**.
- Collaborated in an **Agile** team, utilizing **Git** for source code version control.

PROJECTS

Health Chatbot for Childcare Support

- Built an advanced RAG chatbot using **LlamaIndex**, **Groq API**, and **ChromaDB** (Vector Database) to optimize document embedding storage and semantic search retrieval.
- Processed unstructured data (PDFs/Web) via **BeautifulSoup** and deployed a production-ready interface using **Streamlit** and **FastAPI**.
- **Source:** [Code](#)

Personalized Morning Research Agent

- Architected a **LangGraph** state-machine agent integrated with a **PostgreSQL** database to manage user profile states and session persistence, automating end-to-end research pipelines.
- Integrated **Groq API** for rapid inference alongside **DuckDuckGo API** for real-time aggregation; built an automated multi-format email delivery pipeline via Gmail SMTP, containerized with **Docker Compose**.
- **Source:** [Code](#)

CERTIFICATIONS

IELTS - IDP

September 2024 - September 2026